

Q1

May 17, 2023

```
[63]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix , \
    accuracy_score
```


```
[2]: df = pd.read_csv("AI_heart.csv")
df.head()
```

```
[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52.0	M	0	125	212	0	1	168	0	1.0	2	
1	53.0	M	0	140	203	1	0	155	1	3.1	0	
2	70.0	M	0	145	174	0	1	125	1	2.6	0	
3	61.0	M	0	148	203	0	1	161	0	0.0	2	
4	62.0	F	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
[7]: plt.hist(df['target'])
plt.show()
```



```
[8]: # Data is not imbalanced
```

```
[10]: corr_matrix = df.corr()
      corr_matrix
```

```
[10]:
```

	age	cp	trestbps	chol	fbs	restecg	\
age	1.000000	-0.081711	0.265283	0.217452	0.113552	-0.135061	
cp	-0.081711	1.000000	0.038177	-0.081641	0.079294	0.043581	
trestbps	0.265283	0.038177	1.000000	0.127977	0.181767	-0.123794	
chol	0.217452	-0.081641	0.127977	1.000000	0.026917	-0.147410	
fbs	0.113552	0.079294	0.181767	0.026917	1.000000	-0.104051	
restecg	-0.135061	0.043581	-0.123794	-0.147410	-0.104051	1.000000	
thalach	-0.398372	0.306839	-0.039264	-0.021772	-0.008866	0.048411	
exang	0.086225	-0.401513	0.061197	0.067382	0.049261	-0.065606	
oldpeak	0.196630	-0.174733	0.187434	0.064880	0.010859	-0.050114	
slope	-0.168715	0.131633	-0.120445	-0.014248	-0.061902	0.086086	
ca	0.262819	-0.176206	0.104554	0.074259	0.137156	-0.078072	
thal	0.060679	-0.163341	0.059276	0.100244	-0.042177	-0.020504	
target	-0.222158	0.434854	-0.138772	-0.099966	-0.041164	0.134468	


	thalach	exang	oldpeak	slope	ca	thal	target
age	-0.398372	0.086225	0.196630	-0.168715	0.262819	0.060679	-0.222158

cp	0.306839	-0.401513	-0.174733	0.131633	-0.176206	-0.163341	0.434854
trestbps	-0.039264	0.061197	0.187434	-0.120445	0.104554	0.059276	-0.138772
chol	-0.021772	0.067382	0.064880	-0.014248	0.074259	0.100244	-0.099966
fbs	-0.008866	0.049261	0.010859	-0.061902	0.137156	-0.042177	-0.041164
restecg	0.048411	-0.065606	-0.050114	0.086086	-0.078072	-0.020504	0.134468
thalach	1.000000	-0.380281	-0.349796	0.395308	-0.207888	-0.098068	0.422895
exang	-0.380281	1.000000	0.310844	-0.267335	0.107849	0.197201	-0.438029
oldpeak	-0.349796	0.310844	1.000000	-0.575189	0.221816	0.202672	-0.438441
slope	0.395308	-0.267335	-0.575189	1.000000	-0.073440	-0.094090	0.345512
ca	-0.207888	0.107849	0.221816	-0.073440	1.000000	0.149014	-0.382085
thal	-0.098068	0.197201	0.202672	-0.094090	0.149014	1.000000	-0.337838
target	0.422895	-0.438029	-0.438441	0.345512	-0.382085	-0.337838	1.000000

```
[11]: df.isna().sum()
```

```
[11]: age          102
sex              0
cp              0
trestbps        0
chol            0
fbs            0
restecg         0
thalach         0
exang           0
oldpeak         0
slope           0
ca              0
thal           0
target          0
dtype: int64
```

```
[14]: plt.hist(df['age'],bins=10)
plt.show()
```



```
[16]: df['age'].mean()
```

```
[16]: 54.522210184182015
```

```
[17]: # The maximum patients lie between 50-65, so I will fill the missing values
      ↪ with mean
```

```
[23]: df['age'] = df['age'].fillna(df['age'].mean())
```

```
[24]: df.isna().sum()
```

```
[24]: age      0
      sex      0
      cp      0
      trestbps 0
      chol     0
      fbs      0
      restecg  0
      thalach  0
      exang     0
      oldpeak  0
      slope    0
```

```

ca          0
thal        0
target      0
dtype: int64

```

```
[25]: df.head()
```

```

[25]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0  52.0   M   0     125    212   0         1     168     0       1.0      2
1  53.0   M   0     140    203   1         0     155     1       3.1      0
2  70.0   M   0     145    174   0         1     125     1       2.6      0
3  61.0   M   0     148    203   0         1     161     0       0.0      2
4  62.0   F   0     138    294   1         1     106     0       1.9      1

      ca  thal  target
0     2     3        0
1     0     3        0
2     0     3        0
3     1     3        0
4     3     2        0

```

```
[28]: df['sex'].unique()
```

```
[28]: array(['M', 'F'], dtype=object)
```

```
[31]: # Replace Male with 0 and Female with 1
```

```

[29]: df['sex'] = df['sex'].replace('M',0)
      df['sex'] = df['sex'].replace('F',1)

```

```
[30]: df.head()
```

```

[30]:   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0  52.0   0   0     125    212   0         1     168     0       1.0
1  53.0   0   0     140    203   1         0     155     1       3.1
2  70.0   0   0     145    174   0         1     125     1       2.6
3  61.0   0   0     148    203   0         1     161     0       0.0
4  62.0   1   0     138    294   1         1     106     0       1.9

      slope  ca  thal  target
0         2   2     3        0
1         0   0     3        0
2         0   0     3        0
3         2   1     3        0
4         1   3     2        0

```

```
[32]: df['sex'].unique()
```

```
[32]: array([0, 1])
```

```
[40]: df
```

```
[40]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	\
0	52.00000	0	0	125	212	0	1	168	0	
1	53.00000	0	0	140	203	1	0	155	1	
2	70.00000	0	0	145	174	0	1	125	1	
3	61.00000	0	0	148	203	0	1	161	0	
4	62.00000	1	0	138	294	1	1	106	0	
...
1020	59.00000	0	1	140	221	0	1	164	1	
1021	60.00000	0	0	125	258	0	0	141	1	
1022	47.00000	0	0	110	275	0	0	118	1	
1023	50.00000	1	0	110	254	0	0	159	0	
1024	54.52221	0	0	120	188	0	1	113	0	

	oldpeak	slope	ca	thal	target
0	1.0	2	2	3	0
1	3.1	0	0	3	0
2	2.6	0	0	3	0
3	0.0	2	1	3	0
4	1.9	1	3	2	0
...
1020	0.0	2	0	2	1
1021	2.8	1	1	3	0
1022	1.0	1	1	2	0
1023	0.0	2	0	2	1
1024	1.4	1	1	3	0

```
[1025 rows x 14 columns]
```

```
[42]: # In co-relation matrix, no two features had any close relation, noone was more
      ↪ the 0.8, so everyone is contributing
      # in the dataset
```

```
[44]: df.head()
```

```
[44]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52.0	0	0	125	212	0	1	168	0	1.0	
1	53.0	0	0	140	203	1	0	155	1	3.1	
2	70.0	0	0	145	174	0	1	125	1	2.6	
3	61.0	0	0	148	203	0	1	161	0	0.0	
4	62.0	1	0	138	294	1	1	106	0	1.9	

	slope	ca	thal	target
0	2	2	3	0

1	0	0	3	0
2	0	0	3	0
3	2	1	3	0
4	1	3	2	0

```
[54]: corr_matrix = df.corr()
      corr_matrix
```

```
[54]:
```

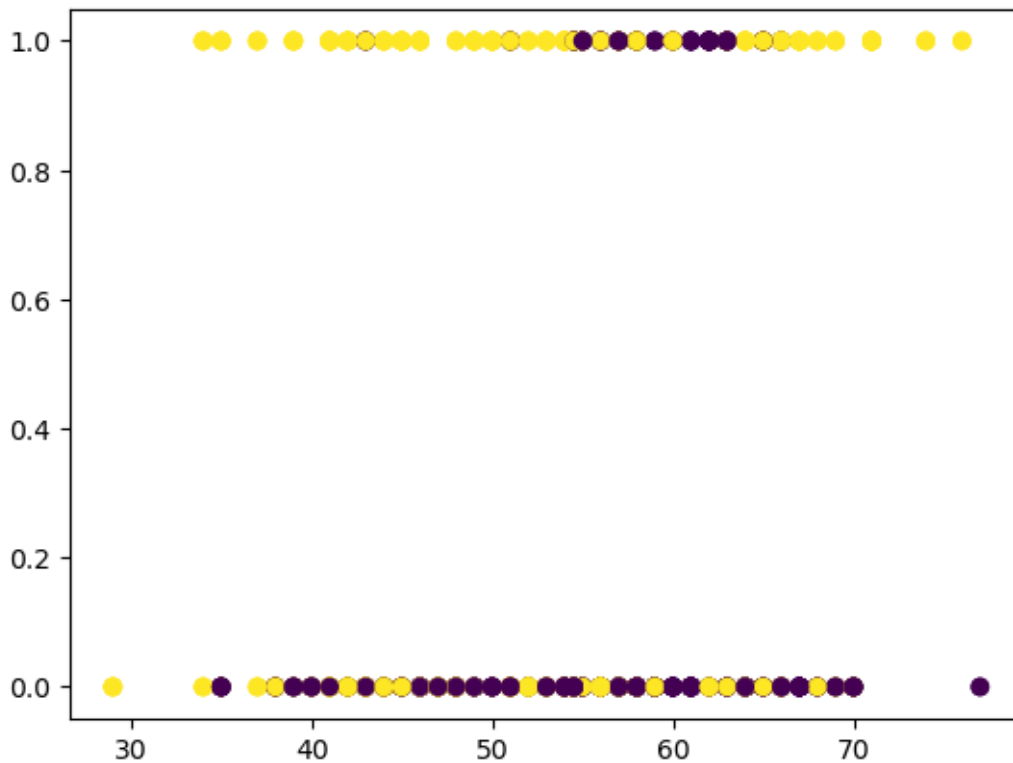
	age	sex	cp	trestbps	chol	fbs	\
age	1.000000	0.102784	-0.077598	0.254872	0.201472	0.106852	
sex	0.102784	1.000000	0.041119	0.078974	0.198258	-0.027200	
cp	-0.077598	0.041119	1.000000	0.038177	-0.081641	0.079294	
trestbps	0.254872	0.078974	0.038177	1.000000	0.127977	0.181767	
chol	0.201472	0.198258	-0.081641	0.127977	1.000000	0.026917	
fbs	0.106852	-0.027200	0.079294	0.181767	0.026917	1.000000	
restecg	-0.127490	0.055117	0.043581	-0.123794	-0.147410	-0.104051	
thalach	-0.376334	0.049365	0.306839	-0.039264	-0.021772	-0.008866	
exang	0.081909	-0.139157	-0.401513	0.061197	0.067382	0.049261	
oldpeak	0.188428	-0.084687	-0.174733	0.187434	0.064880	0.010859	
slope	-0.160123	0.026666	0.131633	-0.120445	-0.014248	-0.061902	
ca	0.252277	-0.111729	-0.176206	0.104554	0.074259	0.137156	
thal	0.057229	-0.198424	-0.163341	0.059276	0.100244	-0.042177	
target	-0.210822	0.279501	0.434854	-0.138772	-0.099966	-0.041164	

	restecg	thalach	exang	oldpeak	slope	ca	\
age	-0.127490	-0.376334	0.081909	0.188428	-0.160123	0.252277	
sex	0.055117	0.049365	-0.139157	-0.084687	0.026666	-0.111729	
cp	0.043581	0.306839	-0.401513	-0.174733	0.131633	-0.176206	
trestbps	-0.123794	-0.039264	0.061197	0.187434	-0.120445	0.104554	
chol	-0.147410	-0.021772	0.067382	0.064880	-0.014248	0.074259	
fbs	-0.104051	-0.008866	0.049261	0.010859	-0.061902	0.137156	
restecg	1.000000	0.048411	-0.065606	-0.050114	0.086086	-0.078072	
thalach	0.048411	1.000000	-0.380281	-0.349796	0.395308	-0.207888	
exang	-0.065606	-0.380281	1.000000	0.310844	-0.267335	0.107849	
oldpeak	-0.050114	-0.349796	0.310844	1.000000	-0.575189	0.221816	
slope	0.086086	0.395308	-0.267335	-0.575189	1.000000	-0.073440	
ca	-0.078072	-0.207888	0.107849	0.221816	-0.073440	1.000000	
thal	-0.020504	-0.098068	0.197201	0.202672	-0.094090	0.149014	
target	0.134468	0.422895	-0.438029	-0.438441	0.345512	-0.382085	

	thal	target
age	0.057229	-0.210822
sex	-0.198424	0.279501
cp	-0.163341	0.434854
trestbps	0.059276	-0.138772
chol	0.100244	-0.099966
fbs	-0.042177	-0.041164

```
restecg -0.020504 0.134468
thalach -0.098068 0.422895
exang 0.197201 -0.438029
oldpeak 0.202672 -0.438441
slope -0.094090 0.345512
ca 0.149014 -0.382085
thal 1.000000 -0.337838
target -0.337838 1.000000
```

```
[56]: plt.scatter(df['age'],df['sex'],c=df['target'])
plt.show()
```



```
[46]: temp = df[df['age']<30]
temp
```

```
[46]:
```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
60	29.0	0	1	130	204	0	0	202	0	0.0	
64	29.0	0	1	130	204	0	0	202	0	0.0	
118	29.0	0	1	130	204	0	0	202	0	0.0	
668	29.0	0	1	130	204	0	0	202	0	0.0	
	slope	ca	thal	target							

60	2	0	2	1
64	2	0	2	1
118	2	0	2	1
668	2	0	2	1


```
[47]: # So we can see yellow means that patient have cancer. Females are more likely
      ↪ to have cancers
```

```
[55]: plt.scatter(df['age'],df['trestbps'],c=df['target'])
```

```
[55]: <matplotlib.collections.PathCollection at 0x7fac694927f0>
```



```
[53]: plt.scatter(df['cp'],df['sex'],c=df['target'])
      plt.show()
```



```
[60]: y = df['target']
X = df.drop('target',axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
[64]: clf = MLPClassifier(hidden_layer_sizes=(100, ), activation='relu',
    ↪solver='adam', max_iter=100)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:",accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.80	0.84	161
1	0.80	0.89	0.84	147
accuracy			0.84	308
macro avg	0.84	0.84	0.84	308
weighted avg	0.85	0.84	0.84	308

Accuracy: 0.8409090909090909

/home/h/anaconda3/lib/python3.9/site-

packages/sklearn/neural_network/_multilayer_perceptron.py:686:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.

warnings.warn(

```
[66]: clf = MLPClassifier(hidden_layer_sizes=(100,100 ), activation='relu',  
    ↪solver='sgd', max_iter=100)  
  
clf.fit(X_train, y_train)  
  
y_pred = clf.predict(X_test)  
  
print(classification_report(y_test, y_pred))  
print("Accuracy:",accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.83	0.28	0.42	161
1	0.54	0.94	0.69	147
accuracy			0.59	308
macro avg	0.69	0.61	0.55	308
weighted avg	0.69	0.59	0.55	308

Accuracy: 0.5941558441558441

/home/h/anaconda3/lib/python3.9/site-

packages/sklearn/neural_network/_multilayer_perceptron.py:686:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.

warnings.warn(

```
[74]: clf = MLPClassifier(hidden_layer_sizes=(10,10 ), activation='relu',  
    ↪solver='adam', max_iter=500)  
  
clf.fit(X_train, y_train)  
  
y_pred = clf.predict(X_test)  
  
print(classification_report(y_test, y_pred))  
print("Accuracy:",accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.78	0.83	161

	1	0.78	0.89	0.83	147
accuracy				0.83	308
macro avg	0.84	0.83	0.83		308
weighted avg	0.84	0.83	0.83		308

Accuracy: 0.8311688311688312

```
[77]: clf = MLPClassifier(hidden_layer_sizes=(20,10,10 ), activation='relu',
    ↪solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

		precision	recall	f1-score	support
	0	0.89	0.78	0.83	161
	1	0.79	0.89	0.84	147
accuracy				0.83	308
macro avg	0.84	0.84	0.83		308
weighted avg	0.84	0.83	0.83		308

Accuracy: 0.8344155844155844

```
[78]: clf = MLPClassifier(hidden_layer_sizes=(20,10,10 ), activation='relu',
    ↪solver='sgd', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

		precision	recall	f1-score	support
	0	0.75	0.71	0.73	161
	1	0.70	0.73	0.72	147
accuracy				0.72	308
macro avg	0.72	0.72	0.72		308
weighted avg	0.73	0.72	0.72		308

Accuracy: 0.724025974025974

```
[83]: clf = MLPClassifier(hidden_layer_sizes=(20,10 ), activation='relu',  
    ↪solver='sgd', max_iter=500)  
  
clf.fit(X_train, y_train)  
  
y_pred = clf.predict(X_test)  
  
print(classification_report(y_test, y_pred))  
print("Accuracy:", accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.52	1.00	0.69	161
1	0.00	0.00	0.00	147
accuracy			0.52	308
macro avg	0.26	0.50	0.34	308
weighted avg	0.27	0.52	0.36	308

Accuracy: 0.5227272727272727

```
/home/h/anaconda3/lib/python3.9/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))  
/home/h/anaconda3/lib/python3.9/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))  
/home/h/anaconda3/lib/python3.9/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))
```

```
[80]: # adam is always giving more accuracy
```

```
[82]: clf = MLPClassifier(hidden_layer_sizes=(20,10 ), activation='relu',  
    ↪solver='adam', max_iter=500)  
  
clf.fit(X_train, y_train)  
  
y_pred = clf.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.82	0.85	161
1	0.82	0.87	0.84	147
accuracy			0.84	308
macro avg	0.84	0.85	0.84	308
weighted avg	0.85	0.84	0.84	308

Accuracy: 0.8441558441558441

```
[88]: clf = MLPClassifier(hidden_layer_sizes=(25,10 ), activation='relu',
    ↪ solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.74	0.80	161
1	0.75	0.87	0.81	147
accuracy			0.80	308
macro avg	0.81	0.80	0.80	308
weighted avg	0.81	0.80	0.80	308

Accuracy: 0.801948051948052

```
[89]: clf = MLPClassifier(solver="lbfgs", hidden_layer_sizes = (20,20), random_state =
    ↪ 1)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.77	0.80	161

	1	0.77	0.84	0.80	147
accuracy				0.80	308
macro avg		0.80	0.80	0.80	308
weighted avg		0.80	0.80	0.80	308

Accuracy: 0.801948051948052

/home/h/anaconda3/lib/python3.9/site-packages/sklearn/neural_network/_multilayer_perceptron.py:541:
 ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)

```
[92]: clf = MLPClassifier(hidden_layer_sizes=(20,20 ), activation='relu',
    ↪solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

		precision	recall	f1-score	support
	0	0.79	0.88	0.84	161
	1	0.85	0.75	0.80	147
accuracy				0.82	308
macro avg		0.82	0.82	0.82	308
weighted avg		0.82	0.82	0.82	308

Accuracy: 0.8181818181818182

```
[93]: clf = MLPClassifier(hidden_layer_sizes=(50,50 ), activation='relu',
    ↪solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.83	0.85	161
1	0.82	0.86	0.84	147
accuracy			0.84	308
macro avg	0.84	0.85	0.84	308
weighted avg	0.85	0.84	0.84	308

Accuracy: 0.8441558441558441

```
[94]: clf = MLPClassifier(hidden_layer_sizes=(100,100 ), activation='relu',
    ↪solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:",accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.77	0.89	0.83	161
1	0.86	0.71	0.78	147
accuracy			0.81	308
macro avg	0.81	0.80	0.80	308
weighted avg	0.81	0.81	0.80	308

Accuracy: 0.8051948051948052

```
[98]: clf = MLPClassifier(hidden_layer_sizes=(50,50 ), activation='relu',
    ↪solver='adam', max_iter=500)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print(classification_report(y_test, y_pred))
print("Accuracy:",accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.78	0.83	161
1	0.79	0.89	0.84	147
accuracy			0.83	308

macro avg	0.84	0.84	0.83	308
weighted avg	0.84	0.83	0.83	308

Accuracy: 0.8344155844155844

```
[99]: # Accuracy at most is 0.84
```

1 Observation

This dataset is small and is not very good to train the whole model. We need more data to get more accurate results

```
[ ]:
```